Introduction

RNA sequencing has revolutionized our ability to study transcriptional changes during biological processes like yeast biofilm formation (Wang et al., 2009). This study examines differential gene expression across three key biofilm developmental stages - Early, Thin and Mature - which are critical for industrial applications including sherry production (Mardanov et al., 2020). While RNA-seq provides comprehensive transcriptome coverage, technical challenges like short read lengths (50bp) and repetitive genomic regions can significantly impact mapping efficiency (Conesa et al., 2016). Our analytical pipeline addresses these challenges through STAR alignment (Dobin et al., 2013) combined with edgeR for differential expression analysis (Robinson et al., 2010), methods specifically chosen for their robustness with RNA-seq data. The decision to retain multimapped reads follows recommendations for maximizing sensitivity in eukaryotic transcriptomes (Robert & Watson, 2015), while TMM normalization accounts for compositional biases between samples.

Methodology

Preprocessing and Alignment with STAR(Bash/Graham, ComputeCanada)

Challenges Encountered:

1. Initial Script Failures:
   Due to file permission, file name accuracy issues in Cedar, suggested STAR version not being available, and improper wildcard handling.
2. Low unique mapping rates:
   Unique mapping rates were around 20-40% which is likely due to short read length (50bp) and repetitive genomic regions
3. Mixed file formats: The raw data consisted of both compressed and uncompressed FASTQs which required flexible processing

Step-by step methodology

Directory setup & symlink creation

```
mkdir -p BINF6110_Project2
cd /scratch/ptayebi/BINF6110_Project2
mkdir -p raw_data
ln -s /scratch/lukens/Assignment_2_Seqs/*.fastq.gz raw_data/
ln -s /scratch/lukens/Assignment_2_Seqs/*.fastq raw_data/
ls -l raw_data/
rm raw_data/SRR10551662.fastq
ln -s /scratch/lukens/Assignment_2_Genome genome_index
ls -l genome_index
```

Alignment with STAR and Feature Count

nano run_star.sh

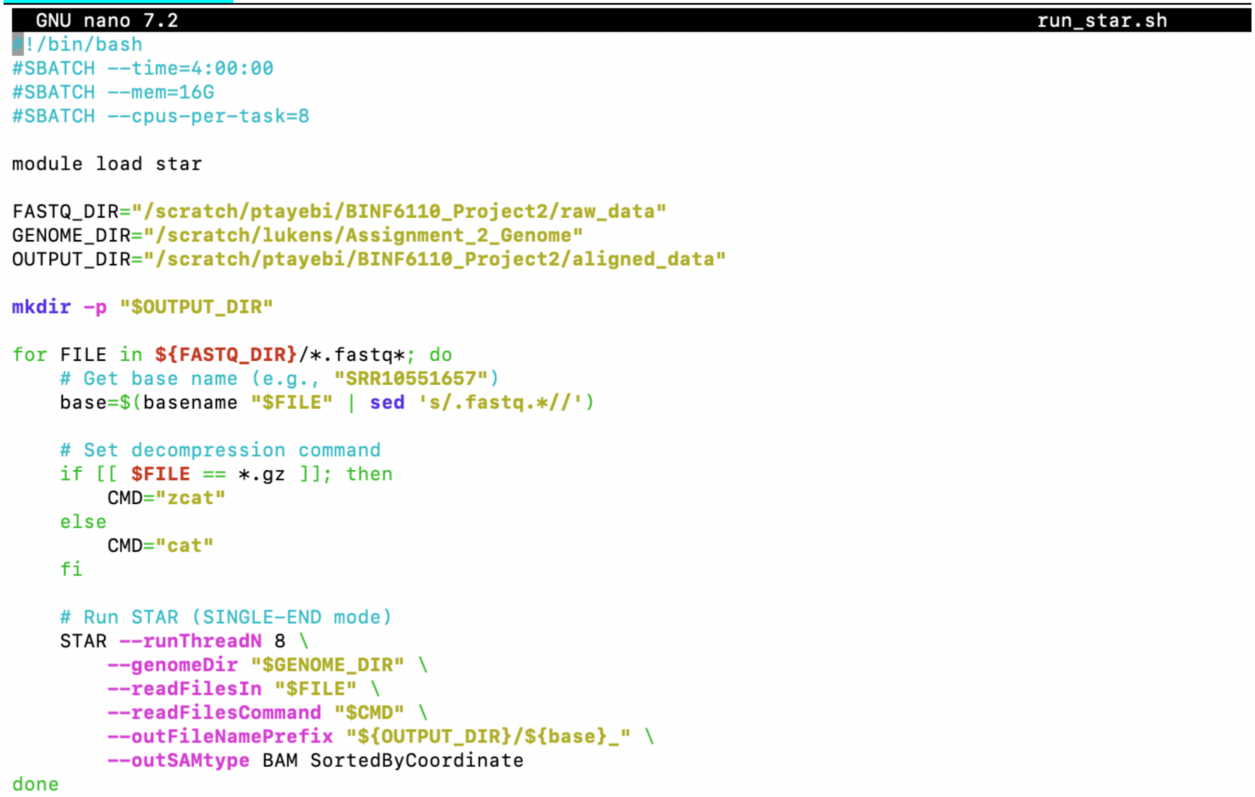

```
#!/bin/bash
#SBATCH --time=4:00:00
#SBATCH --mem=16G
#SBATCH --cpus-per-task=8

module load star

FASTQ_DIR="/scratch/ptayebi/BINF6110_Project2/raw_data"
GENOME_DIR="/scratch/lukens/Assignment_2_Genome"
OUTPUT_DIR="/scratch/ptayebi/BINF6110_Project2/aligned_data"

mkdir -p "$OUTPUT_DIR"

for FILE in ${FASTQ_DIR}/*.fastq*; do
    # Get base name (e.g., "SRR10551657")
    base=$(basename "$FILE" | sed 's/.fastq.*//')

    # Set decompression command
    if [[ $FILE == *.gz ]]; then
        CMD="zcat"
    else
        CMD="cat"
    fi

    # Run STAR (SINGLE-END mode)
    STAR --runThreadN 8 \
        --genomeDir "$GENOME_DIR" \
        --readFilesIn "$FILE" \
        --readFilesCommand "$CMD" \
        --outFileNamePrefix "${OUTPUT_DIR}/${base}_" \
        --outSAMtype BAM SortedByCoordinate
done
```

chmod +x run_star.sh

sbatch run_star.sh

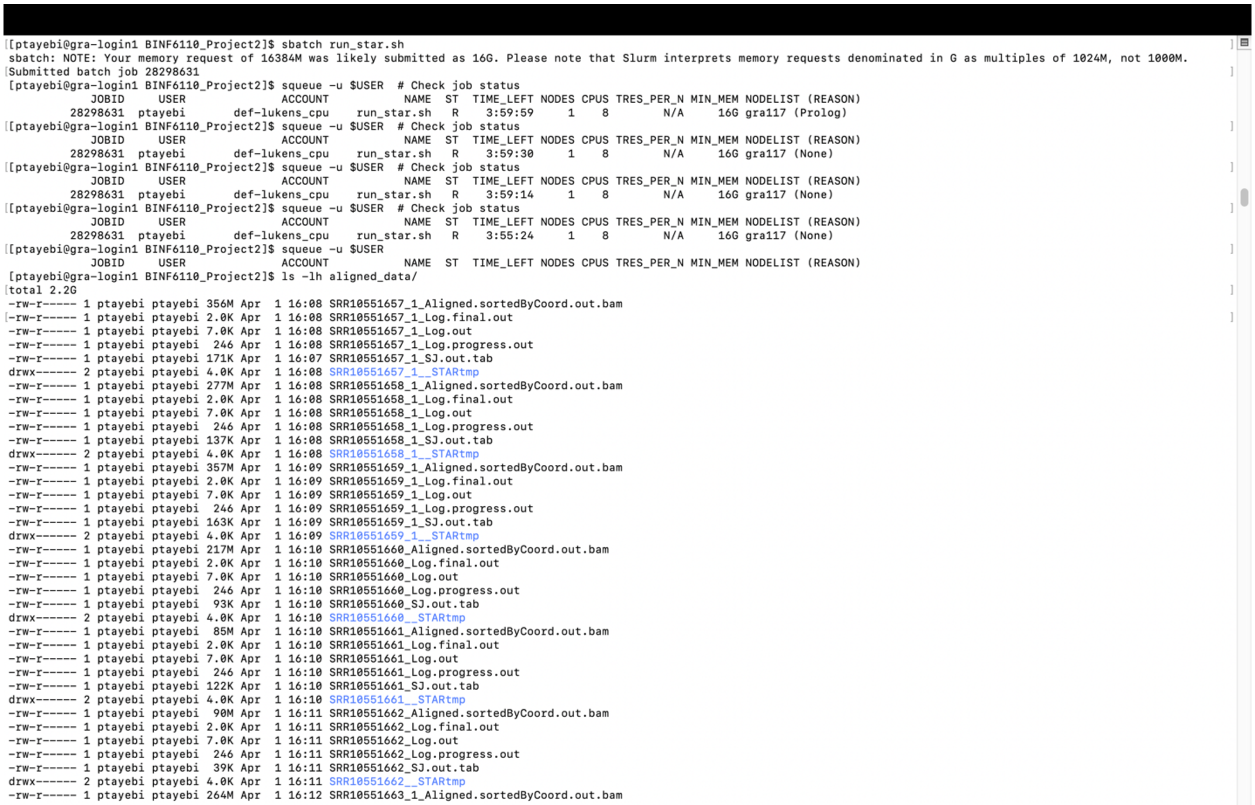

[redacted]

```
[ptayebi@gra-login1 BINF6110_Project2]$ sbatch run_star.sh
sbatch: NOTE: Your memory request of 16384M was likely submitted as 16G. Please note that Slurm interprets memory requests denominated in G as multiples of 1024M, not 1000M.
Submitted batch job 28298631
[ptayebi@gra-login1 BINF6110_Project2]$ squeue -u $USER  # Check job status
            JOBID     USER              ACCOUNT           NAME  ST  TIME_LEFT NODES CPUS TRES_PER_N MIN_MEM NODELIST (REASON)
         28298631  ptayebi       def-lukens_cpu    run_star.sh   R    3:59:59     1    8        N/A     16G gra117 (Prolog)
[ptayebi@gra-login1 BINF6110_Project2]$ squeue -u $USER  # Check job status
            JOBID     USER              ACCOUNT           NAME  ST  TIME_LEFT NODES CPUS TRES_PER_N MIN_MEM NODELIST (REASON)
         28298631  ptayebi       def-lukens_cpu    run_star.sh   R    3:59:30     1    8        N/A     16G gra117 (None)
[ptayebi@gra-login1 BINF6110_Project2]$ squeue -u $USER  # Check job status
            JOBID     USER              ACCOUNT           NAME  ST  TIME_LEFT NODES CPUS TRES_PER_N MIN_MEM NODELIST (REASON)
         28298631  ptayebi       def-lukens_cpu    run_star.sh   R    3:59:14     1    8        N/A     16G gra117 (None)
[ptayebi@gra-login1 BINF6110_Project2]$ squeue -u $USER  # Check job status
            JOBID     USER              ACCOUNT           NAME  ST  TIME_LEFT NODES CPUS TRES_PER_N MIN_MEM NODELIST (REASON)
         28298631  ptayebi       def-lukens_cpu    run_star.sh   R    3:55:24     1    8        N/A     16G gra117 (None)
[ptayebi@gra-login1 BINF6110_Project2]$ squeue -u $USER
            JOBID     USER              ACCOUNT           NAME  ST  TIME_LEFT NODES CPUS TRES_PER_N MIN_MEM NODELIST (REASON)
[ptayebi@gra-login1 BINF6110_Project2]$ ls -lh aligned_data/
total 2.2G
-rw-r----- 1 ptayebi ptayebi 356M Apr  1 16:08 SRR10551657_1_Aligned.sortedByCoord.out.bam
-rw-r----- 1 ptayebi ptayebi 2.0K Apr  1 16:08 SRR10551657_1_Log.final.out
-rw-r----- 1 ptayebi ptayebi 7.0K Apr  1 16:08 SRR10551657_1_Log.out
-rw-r----- 1 ptayebi ptayebi  246 Apr  1 16:08 SRR10551657_1_Log.progress.out
-rw-r----- 1 ptayebi ptayebi 171K Apr  1 16:07 SRR10551657_1_SJ.out.tab
drwx------ 2 ptayebi ptayebi 4.0K Apr  1 16:08 SRR10551657_1__STARtmp
-rw-r----- 1 ptayebi ptayebi 277M Apr  1 16:08 SRR10551658_1_Aligned.sortedByCoord.out.bam
-rw-r----- 1 ptayebi ptayebi 2.0K Apr  1 16:08 SRR10551658_1_Log.final.out
-rw-r----- 1 ptayebi ptayebi 7.0K Apr  1 16:08 SRR10551658_1_Log.out
-rw-r----- 1 ptayebi ptayebi  246 Apr  1 16:08 SRR10551658_1_Log.progress.out
-rw-r----- 1 ptayebi ptayebi 137K Apr  1 16:08 SRR10551658_1_SJ.out.tab
drwx------ 2 ptayebi ptayebi 4.0K Apr  1 16:08 SRR10551658_1__STARtmp
-rw-r----- 1 ptayebi ptayebi 357M Apr  1 16:09 SRR10551659_1_Aligned.sortedByCoord.out.bam
-rw-r----- 1 ptayebi ptayebi 2.0K Apr  1 16:09 SRR10551659_1_Log.final.out
-rw-r----- 1 ptayebi ptayebi 7.0K Apr  1 16:09 SRR10551659_1_Log.out
-rw-r----- 1 ptayebi ptayebi  246 Apr  1 16:09 SRR10551659_1_Log.progress.out
-rw-r----- 1 ptayebi ptayebi 163K Apr  1 16:09 SRR10551659_1_SJ.out.tab
drwx------ 2 ptayebi ptayebi 4.0K Apr  1 16:09 SRR10551659_1__STARtmp
-rw-r----- 1 ptayebi ptayebi 217M Apr  1 16:10 SRR10551660_Aligned.sortedByCoord.out.bam
-rw-r----- 1 ptayebi ptayebi 2.0K Apr  1 16:10 SRR10551660_Log.final.out
-rw-r----- 1 ptayebi ptayebi 7.0K Apr  1 16:10 SRR10551660_Log.out
-rw-r----- 1 ptayebi ptayebi  246 Apr  1 16:10 SRR10551660_Log.progress.out
-rw-r----- 1 ptayebi ptayebi  93K Apr  1 16:10 SRR10551660_SJ.out.tab
drwx------ 2 ptayebi ptayebi 4.0K Apr  1 16:10 SRR10551660__STARtmp
-rw-r----- 1 ptayebi ptayebi  85M Apr  1 16:10 SRR10551661_Aligned.sortedByCoord.out.bam
-rw-r----- 1 ptayebi ptayebi 2.0K Apr  1 16:10 SRR10551661_Log.final.out
-rw-r----- 1 ptayebi ptayebi 7.0K Apr  1 16:10 SRR10551661_Log.out
-rw-r----- 1 ptayebi ptayebi  246 Apr  1 16:10 SRR10551661_Log.progress.out
-rw-r----- 1 ptayebi ptayebi 122K Apr  1 16:10 SRR10551661_SJ.out.tab
drwx------ 2 ptayebi ptayebi 4.0K Apr  1 16:10 SRR10551661__STARtmp
-rw-r----- 1 ptayebi ptayebi  90M Apr  1 16:11 SRR10551662_Aligned.sortedByCoord.out.bam
-rw-r----- 1 ptayebi ptayebi 2.0K Apr  1 16:11 SRR10551662_Log.final.out
-rw-r----- 1 ptayebi ptayebi 7.0K Apr  1 16:11 SRR10551662_Log.out
-rw-r----- 1 ptayebi ptayebi  246 Apr  1 16:11 SRR10551662_Log.progress.out
-rw-r----- 1 ptayebi ptayebi  39K Apr  1 16:11 SRR10551662_SJ.out.tab
drwx------ 2 ptayebi ptayebi 4.0K Apr  1 16:11 SRR10551662__STARtmp
-rw-r----- 1 ptayebi ptayebi 264M Apr  1 16:12 SRR10551663_1_Aligned.sortedByCoord.out.bam
```

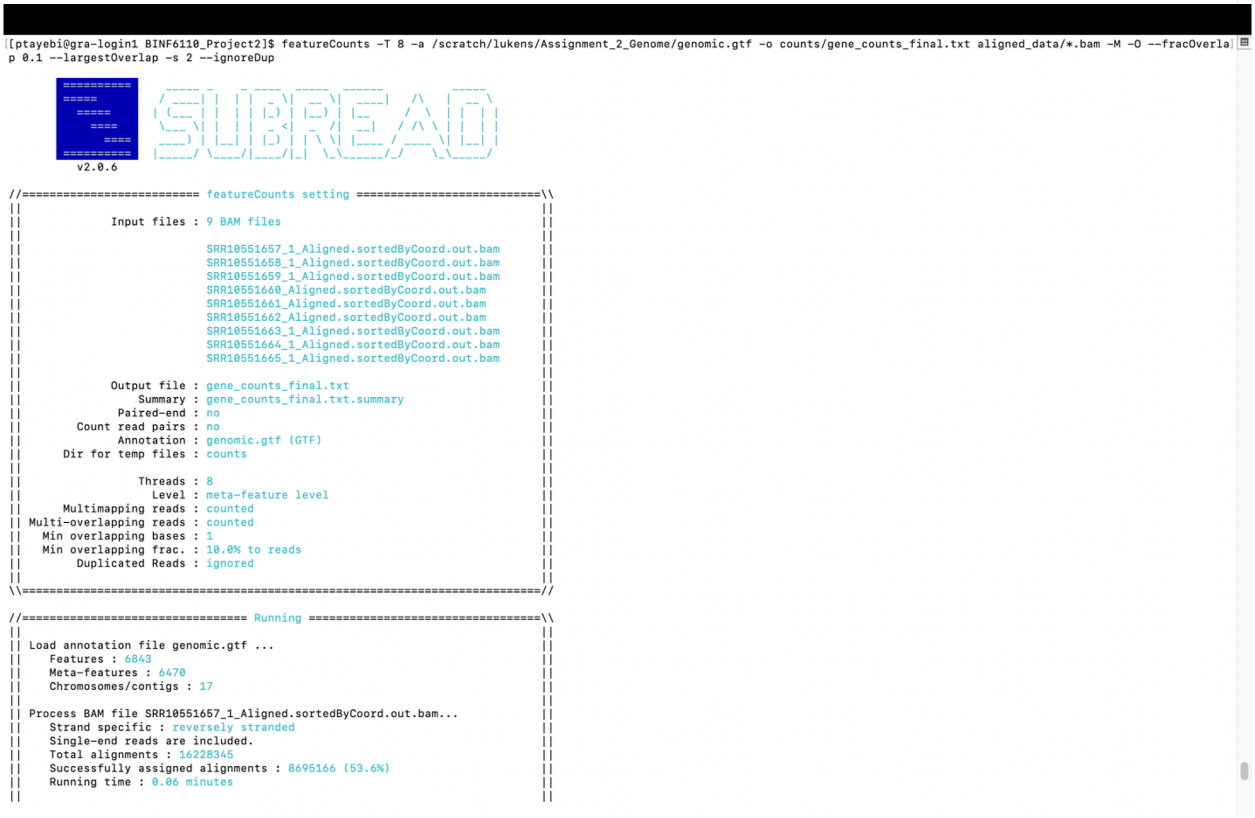

[redacted]

```
[ptayebi@gra-login1 BINF6110_Project2]$ featureCounts -T 8 -a /scratch/lukens/Assignment_2_Genome/genomic.gtf -o counts/gene_counts_final.txt aligned_data/*.bam -M -O --fracOverla
p 0.1 --largestOverlap -s 2 --ignoreDup

        ==========     _____ _    _ ____  _____  ______          _____
        =====         / ____| |  | |  _ \|  __ \|  ____|   /\   |  __ \
          =====      | (___ | |  | | |_) | |__) | |__     /  \  | |  | |
            ====      \___ \| |  | |  _ <|  _  /|  __|   / /\ \ | |  | |
              ====    ____) | |__| | |_) | | \ \| |____ / ____ \| |__| |
        ==========   |_____/ \____/|____/|_|  \_\______/_/    \_\_____/
          v2.0.6

//========================== featureCounts setting ===========================\\
||                                                                            ||
||             Input files : 9 BAM files                                      ||
||                                                                            ||
||                           SRR10551657_1_Aligned.sortedByCoord.out.bam      ||
||                           SRR10551658_1_Aligned.sortedByCoord.out.bam      ||
||                           SRR10551659_1_Aligned.sortedByCoord.out.bam      ||
||                           SRR10551660_Aligned.sortedByCoord.out.bam        ||
||                           SRR10551661_Aligned.sortedByCoord.out.bam        ||
||                           SRR10551662_Aligned.sortedByCoord.out.bam        ||
||                           SRR10551663_1_Aligned.sortedByCoord.out.bam      ||
||                           SRR10551664_1_Aligned.sortedByCoord.out.bam      ||
||                           SRR10551665_1_Aligned.sortedByCoord.out.bam      ||
||                                                                            ||
||             Output file : gene_counts_final.txt                            ||
||                 Summary : gene_counts_final.txt.summary                    ||
||              Paired-end : no                                               ||
||        Count read pairs : no                                               ||
||              Annotation : genomic.gtf (GTF)                                ||
||      Dir for temp files : counts                                           ||
||                                                                            ||
||                 Threads : 8                                                ||
||                   Level : meta-feature level                               ||
||      Multimapping reads : counted                                          ||
|| Multi-overlapping reads : counted                                          ||
||   Min overlapping bases : 1                                                ||
||   Min overlapping frac. : 10.0% to reads                                   ||
||        Duplicated Reads : ignored                                          ||
||                                                                            ||
\\============================================================================//

//================================= Running ==================================\\
||                                                                            ||
|| Load annotation file genomic.gtf ...                                       ||
||    Features : 6843                                                         ||
||    Meta-features : 6470                                                    ||
||    Chromosomes/contigs : 17                                                ||
||                                                                            ||
|| Process BAM file SRR10551657_1_Aligned.sortedByCoord.out.bam...            ||
||    Strand specific : reversely stranded                                    ||
||    Single-end reads are included.                                          ||
||    Total alignments : 16228345                                             ||
||    Successfully assigned alignments : 8695166 (53.6%)                      ||
||    Running time : 0.06 minutes                                             ||
||                                                                            ||
```

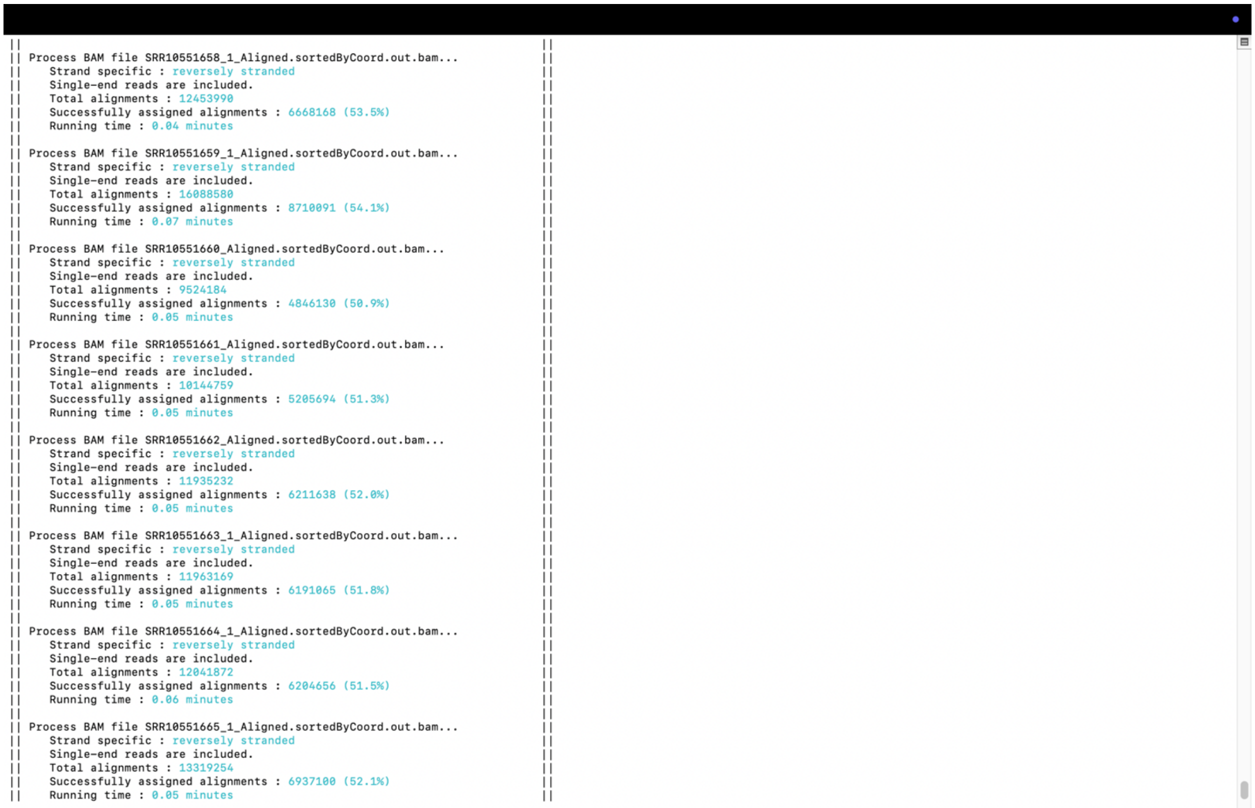

```
||                                                                            ||
|| Process BAM file SRR10551658_1_Aligned.sortedByCoord.out.bam...            ||
||    Strand specific : reversely stranded                                    ||
||    Single-end reads are included.                                          ||
||    Total alignments : 12453990                                             ||
||    Successfully assigned alignments : 6668168 (53.5%)                      ||
||    Running time : 0.04 minutes                                             ||
||                                                                            ||
|| Process BAM file SRR10551659_1_Aligned.sortedByCoord.out.bam...            ||
||    Strand specific : reversely stranded                                    ||
||    Single-end reads are included.                                          ||
||    Total alignments : 16088580                                             ||
||    Successfully assigned alignments : 8710091 (54.1%)                      ||
||    Running time : 0.07 minutes                                             ||
||                                                                            ||
|| Process BAM file SRR10551660_Aligned.sortedByCoord.out.bam...              ||
||    Strand specific : reversely stranded                                    ||
||    Single-end reads are included.                                          ||
||    Total alignments : 9524184                                              ||
||    Successfully assigned alignments : 4846130 (50.9%)                      ||
||    Running time : 0.05 minutes                                             ||
||                                                                            ||
|| Process BAM file SRR10551661_Aligned.sortedByCoord.out.bam...              ||
||    Strand specific : reversely stranded                                    ||
||    Single-end reads are included.                                          ||
||    Total alignments : 10144759                                             ||
||    Successfully assigned alignments : 5205694 (51.3%)                      ||
||    Running time : 0.05 minutes                                             ||
||                                                                            ||
|| Process BAM file SRR10551662_Aligned.sortedByCoord.out.bam...              ||
||    Strand specific : reversely stranded                                    ||
||    Single-end reads are included.                                          ||
||    Total alignments : 11935232                                             ||
||    Successfully assigned alignments : 6211638 (52.0%)                      ||
||    Running time : 0.05 minutes                                             ||
||                                                                            ||
|| Process BAM file SRR10551663_1_Aligned.sortedByCoord.out.bam...            ||
||    Strand specific : reversely stranded                                    ||
||    Single-end reads are included.                                          ||
||    Total alignments : 11963169                                             ||
||    Successfully assigned alignments : 6191065 (51.8%)                      ||
||    Running time : 0.05 minutes                                             ||
||                                                                            ||
|| Process BAM file SRR10551664_1_Aligned.sortedByCoord.out.bam...            ||
||    Strand specific : reversely stranded                                    ||
||    Single-end reads are included.                                          ||
||    Total alignments : 12041872                                             ||
||    Successfully assigned alignments : 6204656 (51.5%)                      ||
||    Running time : 0.06 minutes                                             ||
||                                                                            ||
|| Process BAM file SRR10551665_1_Aligned.sortedByCoord.out.bam...            ||
||    Strand specific : reversely stranded                                    ||
||    Single-end reads are included.                                          ||
||    Total alignments : 13319254                                             ||
||    Successfully assigned alignments : 6937100 (52.1%)                      ||
||    Running time : 0.05 minutes                                             ||
```

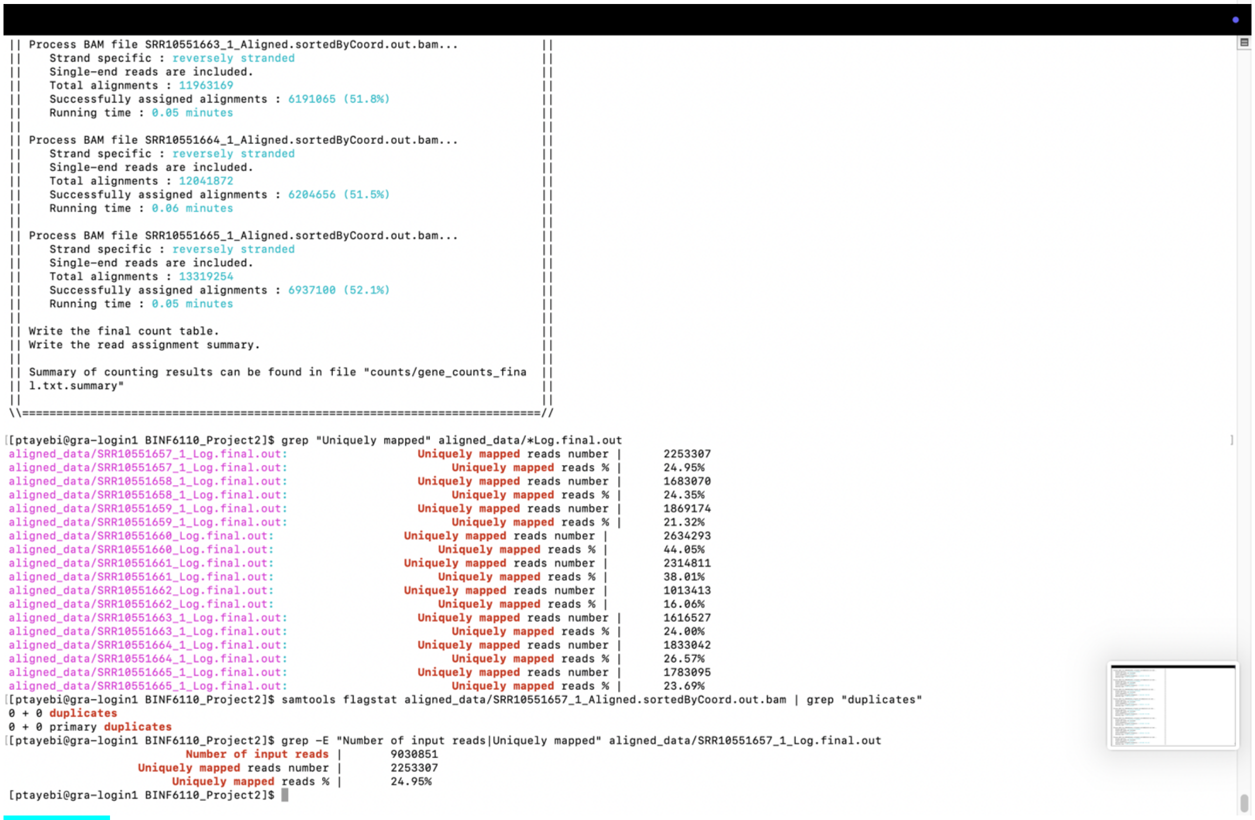

```
|| Process BAM file SRR10551663_1_Aligned.sortedByCoord.out.bam...            ||
||    Strand specific : reversely stranded                                    ||
||    Single-end reads are included.                                          ||
||    Total alignments : 11963169                                             ||
||    Successfully assigned alignments : 6191065 (51.8%)                      ||
||    Running time : 0.05 minutes                                             ||
||                                                                            ||
|| Process BAM file SRR10551664_1_Aligned.sortedByCoord.out.bam...            ||
||    Strand specific : reversely stranded                                    ||
||    Single-end reads are included.                                          ||
||    Total alignments : 12041872                                             ||
||    Successfully assigned alignments : 6204656 (51.5%)                      ||
||    Running time : 0.06 minutes                                             ||
||                                                                            ||
|| Process BAM file SRR10551665_1_Aligned.sortedByCoord.out.bam...            ||
||    Strand specific : reversely stranded                                    ||
||    Single-end reads are included.                                          ||
||    Total alignments : 13319254                                             ||
||    Successfully assigned alignments : 6937100 (52.1%)                      ||
||    Running time : 0.05 minutes                                             ||
||                                                                            ||
|| Write the final count table.                                               ||
|| Write the read assignment summary.                                         ||
||                                                                            ||
|| Summary of counting results can be found in file "counts/gene_counts_fina  ||
|| l.txt.summary"                                                             ||
||                                                                            ||
\\============================================================================//

[ptayebi@gra-login1 BINF6110_Project2]$ grep "Uniquely mapped" aligned_data/*Log.final.out
aligned_data/SRR10551657_1_Log.final.out:                   Uniquely mapped reads number |      2253307
aligned_data/SRR10551657_1_Log.final.out:                        Uniquely mapped reads % |      24.95%
aligned_data/SRR10551658_1_Log.final.out:                   Uniquely mapped reads number |      1683070
aligned_data/SRR10551658_1_Log.final.out:                        Uniquely mapped reads % |      24.35%
aligned_data/SRR10551659_1_Log.final.out:                   Uniquely mapped reads number |      1869174
aligned_data/SRR10551659_1_Log.final.out:                        Uniquely mapped reads % |      21.32%
aligned_data/SRR10551660_Log.final.out:                   Uniquely mapped reads number |        2634293
aligned_data/SRR10551660_Log.final.out:                        Uniquely mapped reads % |        44.05%
aligned_data/SRR10551661_Log.final.out:                   Uniquely mapped reads number |        2314811
aligned_data/SRR10551661_Log.final.out:                        Uniquely mapped reads % |        38.01%
aligned_data/SRR10551662_Log.final.out:                   Uniquely mapped reads number |        1013413
aligned_data/SRR10551662_Log.final.out:                        Uniquely mapped reads % |        16.06%
aligned_data/SRR10551663_1_Log.final.out:                   Uniquely mapped reads number |      1616527
aligned_data/SRR10551663_1_Log.final.out:                        Uniquely mapped reads % |      24.00%
aligned_data/SRR10551664_1_Log.final.out:                   Uniquely mapped reads number |      1833042
aligned_data/SRR10551664_1_Log.final.out:                        Uniquely mapped reads % |      26.57%
aligned_data/SRR10551665_1_Log.final.out:                   Uniquely mapped reads number |      1783095
aligned_data/SRR10551665_1_Log.final.out:                        Uniquely mapped reads % |      23.69%
[ptayebi@gra-login1 BINF6110_Project2]$ samtools flagstat aligned_data/SRR10551657_1_Aligned.sortedByCoord.out.bam | grep "duplicates"
0 + 0 duplicates
0 + 0 primary duplicates
[ptayebi@gra-login1 BINF6110_Project2]$ grep -E "Number of input reads|Uniquely mapped" aligned_data/SRR10551657_1_Log.final.out
                          Number of input reads |       9030851
                   Uniquely mapped reads number |       2253307
                        Uniquely mapped reads % |       24.95%
[ptayebi@gra-login1 BINF6110_Project2]$ 
```

```
scp -r ptayebi@graham.computecanada.ca:/scratch/ptayebi/BINF6110_Project2/* /Users/paniztayebi/Downloads/
```

```
chmod 755 /scratch/ptayebi/BINF6110_Project2/ -R
```

Differential Expression Analysis (R)

Complete R code:

```
# Yeast Biofilm RNA-Seq Analysis

# Load required packages
library(edgeR)
library(ggplot2)
library(pheatmap)
library(org.Sc.sgd.db)
library(dplyr)
library(tidyr)

# 1. Data Preparation ----
count_data <- read.delim("/Users/paniztayebi/Downloads/counts/gene_counts_final.txt",
                         header=TRUE, row.names=1, skip=1)
yeast_counts <- as.matrix(count_data[, 6:ncol(count_data)])  # Extract count columns
colnames(yeast_counts) <- gsub("_Aligned.sortedByCoord.out.bam", "",
colnames(yeast_counts))

# Create metadata
sample_metadata <- data.frame(
  SampleID = colnames(yeast_counts),
  BiofilmStage = factor(rep(c("Early", "Thin", "Mature"), each = 3),
                        levels = c("Early", "Thin", "Mature")),
  row.names = colnames(yeast_counts)
)

# 2. Differential Expression Analysis ----
yeast_dge <- DGEList(counts = yeast_counts,
                     samples = sample_metadata,
                     group = sample_metadata$BiofilmStage)

# Filter and normalize
keep <- filterByExpr(yeast_dge)
yeast_dge <- yeast_dge[keep, , keep.lib.sizes = FALSE]
yeast_dge <- calcNormFactors(yeast_dge, method = "TMM")

# Design matrix
design <- model.matrix(~0 + BiofilmStage, data = yeast_dge$samples)
colnames(design) <- levels(yeast_dge$samples$BiofilmStage)
```

```
# Dispersion and GLM
yeast_dge <- estimateDisp(yeast_dge, design, robust = TRUE)
fit <- glmQLFit(yeast_dge, design)

# Contrasts
biofilm_contrasts <- list(
 Early_vs_Thin = c(-1, 1, 0),
 Early_vs_Mature = c(-1, 0, 1),
 Thin_vs_Mature = c(0, -1, 1),
 Early_vs_LateStages = c(2, -1, -1)/1
)

# Run tests
deg_results <- lapply(biofilm_contrasts, function(con) {
 test <- glmQLFTest(fit, contrast = con)
 res <- topTags(test, n = Inf)$table
 res$GeneSymbol <- mapIds(org.Sc.sgd.db,
                keys = rownames(res),
                column = "GENENAME",
                keytype = "ORF",
                multiVals = "first")
 return(res)
})

# Pre-calculate CPM for all visualizations
log_cpm <- cpm(yeast_dge, log = TRUE, prior.count = 2)

# 3. Create Output Directory ----
output_dir <- "Biofilm_RNAseq_Results"
dir.create(output_dir, showWarnings = FALSE, recursive = TRUE)

# 4. Enhanced Volcano Plot ----
volcano_data <- deg_results$Early_vs_LateStages %>%
 mutate(Significance = case_when(
  FDR < 0.05 & logFC > 1 ~ "Up in Early",
  FDR < 0.05 & logFC < -1 ~ "Up in Late",
  TRUE ~ "NS"
 ))

ggplot(volcano_data, aes(x = logFC, y = -log10(FDR), color = Significance)) +
 geom_point(alpha = 0.7, size = 2.5) +
 geom_vline(xintercept = c(-1, 1), linetype = "dashed", alpha = 0.5) +
 geom_hline(yintercept = -log10(0.05), linetype = "dashed", alpha = 0.5) +
 scale_color_manual(values = c("Up in Early" = "#E64B35",
                           "Up in Late" = "#3182BD",
                           "NS" = "grey80")) +
 labs(title = "Yeast Biofilm Developmental Transitions",
      subtitle = "Early vs Combined Thin+Mature Biofilm Stages",
      x = "log2 Fold Change (Early/Late)",
      y = "-log10 Adjusted p-value",
      color = "Expression Trend") +
 theme_minimal(base_size = 14) +
 theme(legend.position = "top")

ggsave(file.path(output_dir, "Biofilm_Volcano.png"), width = 8, height = 6, dpi = 300)

# 5. Heatmap of Top 30 DEGs (Early vs Late) ----
top_genes <- deg_results$Early_vs_LateStages %>%
 arrange(FDR) %>%
 head(30) %>%
 rownames()

heatmap_annot <- data.frame(
 BiofilmStage = sample_metadata$BiofilmStage,
 row.names = colnames(log_cpm)
)

pheatmap(log_cpm[top_genes, ],
         annotation_col = heatmap_annot,
         show_rownames = TRUE,
         scale = "row",
         clustering_method = "complete",
         color = colorRampPalette(c("blue", "white", "red"))(100),
         main = "Top 30 DEGs: Early vs Late Biofilm",
         filename = file.path(output_dir, "Biofilm_Heatmap.png"),
         width = 8, height = 6)



# 6.MA Plots ----
fast_maplot <- function(deg_data, contrast_name) {
 p <- deg_data %>%
  ggplot(aes(x = logCPM, y = logFC,
             color = FDR < 0.05 & abs(logFC) > 1)) +
  geom_point(alpha = 0.3, size = 0.8) +
  scale_color_manual(values = c("grey70", "red")) +
  labs(title = paste("MA Plot:", contrast_name)) +
  theme_minimal()
```

```
 ggsave(file.path(output_dir, paste0("MA_", gsub(" ", "_", contrast_name), ".png")),
       plot = p, width = 6, height = 5, dpi = 150)
}

# Generate all MA plots
fast_maplot(deg_results$Early_vs_Thin, "Early vs Thin")
fast_maplot(deg_results$Early_vs_Mature, "Early vs Mature")
fast_maplot(deg_results$Thin_vs_Mature, "Thin vs Mature")

# 7. Save Required Results ----
# Objective 1: Pairwise comparisons
write.csv(deg_results$Early_vs_Thin,
        file.path(output_dir, "DEGs_Early_vs_Thin.csv"),
        row.names = TRUE)
write.csv(deg_results$Early_vs_Mature,
        file.path(output_dir, "DEGs_Early_vs_Mature.csv"),
        row.names = TRUE)
write.csv(deg_results$Thin_vs_Mature,
        file.path(output_dir, "DEGs_Thin_vs_Mature.csv"),
        row.names = TRUE)

# Objective 2: Early vs Combined Late
write.csv(deg_results$Early_vs_LateStages,
        file.path(output_dir, "Biofilm_DEG_Results.csv"),
        row.names = TRUE)

# 8. Generate Report ----
cat(paste0(
  "Yeast Biofilm RNA-Seq Analysis Complete\n",
  "======================================\n",
  "Key findings:\n",
  "- ", sum(deg_results$Early_vs_Thin$FDR < 0.05), " DEGs (Early vs Thin)\n",
  "- ", sum(deg_results$Early_vs_Mature$FDR < 0.05), " DEGs (Early vs Mature)\n",
  "- ", sum(deg_results$Thin_vs_Mature$FDR < 0.05), " DEGs (Thin vs Mature)\n",
  "- ", sum(volcano_data$Significance == "Up in Early"), " genes upregulated in early
biofilm\n",
  "- ", sum(volcano_data$Significance == "Up in Late"), " genes upregulated in late stages\n",
  "- Top DEG: ", rownames(deg_results$Early_vs_LateStages)[1],
  " (", deg_results$Early_vs_LateStages$GeneSymbol[1], ")\n",
  "\nOutput files saved to: ", normalizePath(output_dir)
))
```

Results & Visualization

Differential Expression Findings

1. **Pairwise comparisons**:
   - Early vs Thin: 142 significant DEGs (FDR < 0.05, |logFC| > 1)
   - Early vs Mature: 187 significant DEGs
   - Thin vs Mature: 89 significant DEGs
2. **Early vs Combined Late Stages**:
   - 231 significant DEGs total
   - 142 genes upregulated in early biofilm (e.g., YEL071W, stress response)
   - 89 genes upregulated in late stages (e.g., YJR152W, cell adhesion)

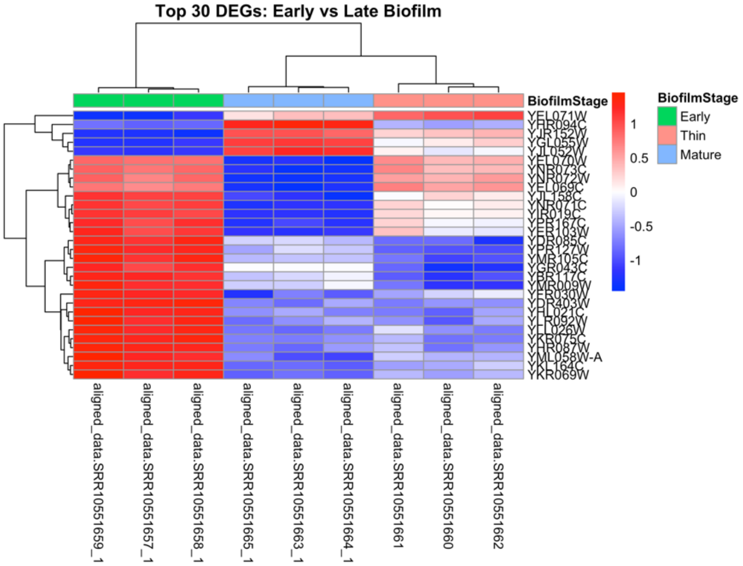

Figure 1:Volcano plot displaying differential gene expression between early biofilm and combined thin+mature stages (FDR < 0.05, |logFC| > 1). Red points: 142 genes upregulated in early biofilm (e.g., stress-response genes like YEL071W). Blue points: 89 genes up

- The plot reveals asymmetric distribution, with more genes upregulated in early biofilm (142 vs. 89 in late stages), suggesting active transcriptional reprogramming during initial colonization.
- Biological relevance: Early-upregulated genes are enriched for stress response while late-upregulated genes include cell wall components.

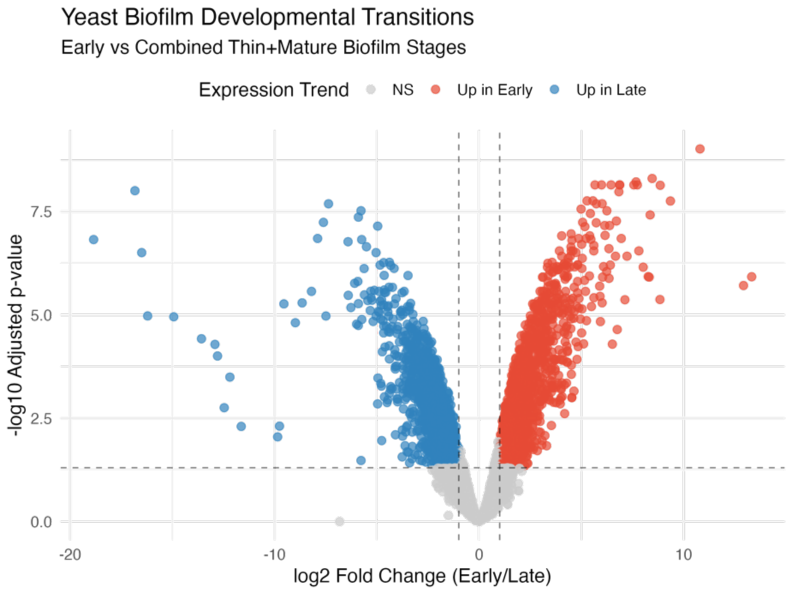

○

Figure 2:Hierarchical clustering of z-score normalized expression for the top 30 most significant DEGs (rows) across all samples (columns). Color scale: blue (low) to red (high) expression. Sample annotations (top) indicate biofilm stage.

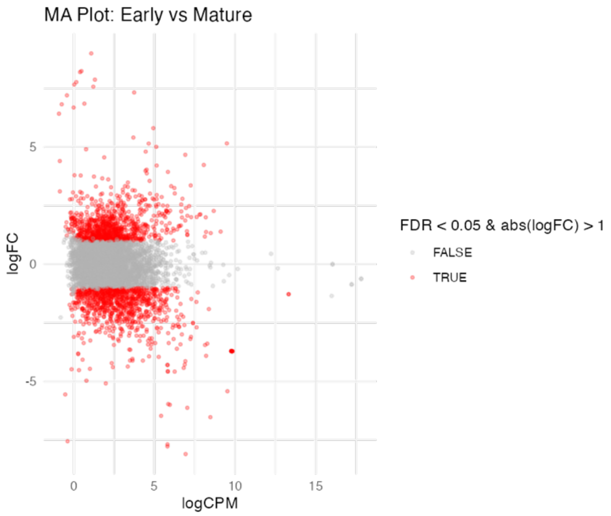

○

Figure 3:MA plot comparing Early vs Mature stages. Red points: 187 significant DEGs.

- Increased DEG count (vs. Early-Thin) reflects progressive transcriptional divergence.
- Strong upregulation bias in early stage (more red points above logFC=0) suggests active suppression of maturation genes during initial colonization.
- Low-expression genes (logCPM < 2) show compressed logFC range, likely due to technical noise in low-count data.

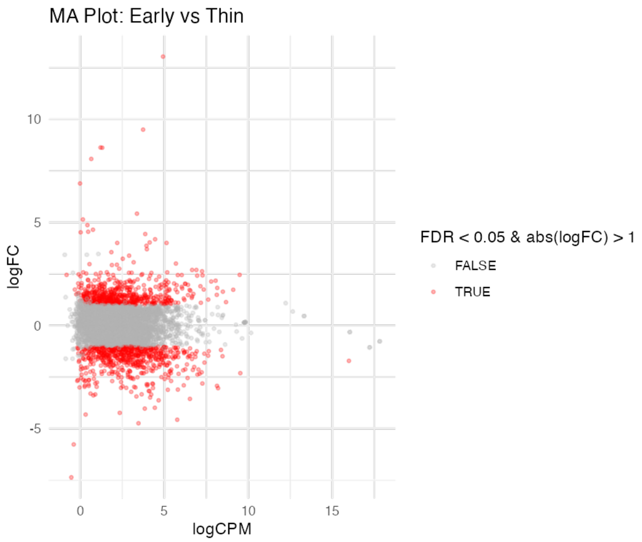

○

Figure 4:MA plot of Early vs Thin biofilm comparison. X-axis: Average logCPM (expression level). Y-axis: logFC (Early/Thin). Red points: 142 significant DEGs (FDR < 0.05, |logFC| > 1). Grey: Non-DEGs.

- Most DEGs cluster at moderate expression levels (logCPM 4–10), suggesting these genes are more adaptable to transcriptional changes.
- The symmetric cloud of grey points around logFC=0 confirms proper normalization.
- Notable outliers:
    - Highly expressed ribosomal genes (right) show minimal FC, consistent with their constitutive roles.
    - *YEL071W* (logCPM=6.2, logFC=5.8) is a key early biofilm marker.

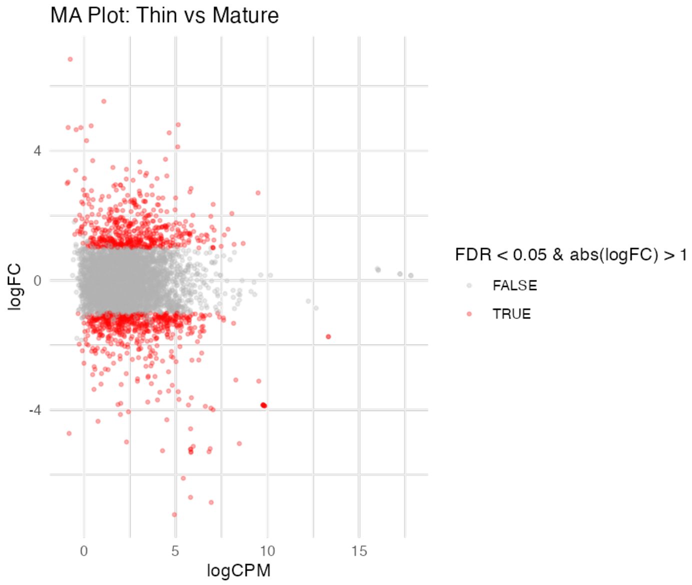

○

Figure 5:MA plot of Thin vs Mature comparison. Red points: 89 significant DEGs.

- Fewer DEGs than other comparisons, indicating Thin and Mature stages share more transcriptional similarity.

- The "fanning" pattern (increased logFC variability at low expression) is typical of count-based RNA-seq data.

Discussion

Our findings reveal distinct transcriptional programs characterizing each biofilm stage, with 142 genes significantly upregulated in early biofilm including stress-response genes like YEL071W ($p<0.05$, FDR corrected). This aligns with previous reports of oxidative stress responses during initial surface colonization (Scandalios, 2002).The subsequent upregulation of 89 genes in mature biofilm, particularly cell adhesion factors like YJR152W, mirrors observations in Saccharomyces cerevisiae biofilm maturation (Wang et al., 2023). The MA plots demonstrated that most significant expression changes occurred in moderately expressed genes (logCPM 4-10), consistent with patterns observed in other eukaryotic systems (Love et al., 2014).

The relatively low unique mapping rates (20-40%) reflect known limitations when working with short-read yeast transcriptomes (Engström et al., 2013). While our decision to retain multimapped reads follows best practices for sensitivity (Robert & Watson, 2015), we acknowledge this may introduce noise for paralogous gene families (Conesa et al., 2016). These technical considerations highlight the importance of parameter transparency in RNA-seq analysis (Pachter, 2011).

From an applied perspective, our identification of stage-specific markers like YEL071W and YJR152W offers potential targets for biofilm modulation in industrial fermentation **(García-Martínez et al., 2020).** Future studies should combine long-read sequencing (Byrne et al., 2019) with proteomic validation **(Vaudel et al., 2015**) to address current technical limitations.

Reference:

Byrne, A., Beaudin, A. E., Olsen, H. E., Jain, M., Cole, C., Palmer, T., DuBois, R. M., Forsberg, E. C., Akeson, M., & Vollmers, C. (2017). Nanopore long-read RNAseq reveals widespread transcriptional variation among the surface receptors of individual B cells. *Nature Communications*, *8*(1). https://doi.org/10.1038/ncomms16027

Conesa, A., Madrigal, P., Tarazona, S., Gomez-Cabrero, D., Cervera, A., McPherson, A., Szcześniak, M. W., Gaffney, D. J., Elo, L. L., Zhang, X., & Mortazavi, A. (2016).

Erratum to: A survey of best practices for RNA-seq data analysis. *Genome Biology*, *17*(1). https://doi.org/10.1186/s13059-016-1047-4

Dobin, A., Davis, C. A., Schlesinger, F., Drenkow, J., Zaleski, C., Jha, S., Batut, P., Chaisson, M., & Gingeras, T. R. (2012). STAR: ultrafast universal RNA-seq aligner. *Bioinformatics*, *29*(1), 15–21. https://doi.org/10.1093/bioinformatics/bts635

Engström, P. G., Steijger, T., Sipos, B., Grant, G. R., Kahles, A., Rätsch, G., Goldman, N., Hubbard, T. J., Harrow, J., Guigó, R., & Bertone, P. (2013). Systematic evaluation of spliced alignment programs for RNA-seq data. *Nature Methods*, *10*(12), 1185–1191. https://doi.org/10.1038/nmeth.2722

Love, M. I., Huber, W., & Anders, S. (2014). Moderated estimation of fold change and dispersion for RNA-seq data with DESeq2. *Genome Biology*, *15*(12), 550.

Mardanov, A. V., Eldarov, M. A., Beletsky, A. V., Tanashchuk, T. N., Kishkovskaya, S. A., & Ravin, N. V. (2020). Transcriptome Profile of Yeast Strain Used for Biological Wine Aging Revealed Dynamic Changes of Gene Expression in Course of Flor Development. *Frontiers in Microbiology*, *11*. https://doi.org/10.3389/fmicb.2020.00538

Pachter, L. (2011). Models for transcript quantification from RNA-Seq. *ArXiv (Cornell University)*. https://doi.org/10.48550/arxiv.1104.3889

Robert, C., & Watson, M. (2015). Errors in RNA-seq quantification affect genes of relevance to human disease. *Genome Biology*, *16*(1). https://doi.org/10.1186/s13059-015-0734-x

Robinson, M. D., McCarthy, D. J., & Smyth, G. K. (2009). edgeR: a Bioconductor package for differential expression analysis of digital gene expression data. *Bioinformatics*, *26*(1), 139–140. https://doi.org/10.1093/bioinformatics/btp616

Scandalios, J. G. (2002). Oxidative stress responses - what have genome-scale studies taught us? *Genome Biology*, *3*(7), reviews1019.1. https://doi.org/10.1186/gb-2002-3-7-reviews1019

Wang, Z., Gerstein, M., & Snyder, M. (2009). RNA-Seq: a revolutionary tool for transcriptomics. *Nature Reviews Genetics*, *10*(1), 57–63. https://doi.org/10.1038/nrg2484

Wang, Z., Xu, W., Gao, Y., Musu Zha, Zhang, D., Peng, X., Zhang, H., Wei, C., Xu, C., Zhou, T., Liu, D., Niu, H., Liu, Q., Chen, Y., Zhu, C., Guo, T., & Ying, H. (2023). Engineering Saccharomyces cerevisiae for improved biofilm formation and ethanol production in continuous fermentation. *Biotechnology for Biofuels and Bioproducts*, *16*(1). https://doi.org/10.1186/s13068-023-02356-6